

covtl-pipeline

User Manual

Frédéric Berthommier

`fberthommier@users.noreply.github.com`

Version 1.0.2 — September 2026

Licensed under GPL-3.0-or-later

Abstract

covtl-pipeline is a Python package and command-line tool (**vtl-synth**) that converts raw English text — or direct SAMPA phonetic input — into an articulatory speech signal and an animated video of the vocal-tract sagittal cross-section. Its synthesis engine implements the **Berthommier COVTL articulatory model**: every phoneme carries a target defined by *polar coordinates* (ρ, θ) in a compact two-dimensional phonetic plane, and coarticulation emerges from the *superposition* of a continuous vocalic trajectory with transient consonantal gestures projected onto disjoint sets of VocalTractLab parameters. A tag-based extension engine drives voicing, nasality, laterality, glottal effort and F0, and a syllabic trajectory engine (*syltraj*) shapes the consonant/vowel arcs, inter-word pauses and phrase-final holds. The 400 Hz tract sequence is synthesized *directly* by this model — there is no intermediate gestural-score stage — and audio is rendered by **VocalTractLab 2.4** (Peter Birkholz, TU Dresden) through the **vocaltractlab-cython** wheel with the JD3 speaker. This manual documents the model in detail (polar targets, projection coefficients, coarticulation rules), the command-line interface, the Python API, the file formats, and the underlying mechanisms and constants.

At a glance

- **Pipeline:** text $\rightarrow$ SAMPA $\rightarrow$ polar targets $(\rho, \theta) \rightarrow$ 100 Hz model trajectory $\rightarrow$ `.tract` (400 Hz, 19 + 11 params) $\rightarrow$ `.wav` (44.1 kHz) $\rightarrow$ `.mp4`
- **CLI:** `vtl-synth run "this is easy for us" -o ./out`
- **Python API:** `from vtl_synth import Pipeline`
- **Model:** Berthommier COVTL — polar phonetic targets, selector-based coarticulation, S-curve arcs
- **Synthesizer:** VocalTractLab 2.4 (JD3 speaker) via **vocaltractlab-cython**
- **License:** GPL-3.0-or-later (inherited from upstream VTL components)

Contents

1	Introduction	3
2	Quick start	4
2.1	Prerequisites	4
2.2	Installation	4
2.3	First run	4
2.4	Inspecting the output	5
3	The COVTL model I: phonetic targets in polar coordinates	5
3.1	The (ρ, θ) phonetic plane	5
3.2	Vowel targets	6
3.3	Consonant targets and selectors	6
3.4	Context-conditioned targets	8
3.5	Projection onto the VocalTractLab parameters	8
4	The COVTL model II: coarticulation	10
4.1	Superposition of the vocalic and consonantal branches	10
4.2	Lightened selectors: the tongue during labials	10
4.3	The transition arc: an asymmetric S-curve in the plane	11
4.4	Clusters: per-parameter scoping	11
4.5	Vowel-to-vowel coarticulation and the schwa waypoint	11
4.6	Coarticulation across word boundaries	12
4.7	Nasal carry-over	12
5	The syllabic trajectory engine (syltraj)	12
5.1	The orthogonal branch	12
5.2	The legacy engine	13
6	The source model: envelopes and the glottal source	13
6.1	The three-level tag system	13
6.2	The amplitude envelope $E(t)$ and the source timetable	13
6.3	Voicing, VOT and nasality	14
6.4	Glottal source, pressure and effort	14
6.5	F0: declination and micro-prosody	14
7	Architecture and data flow	15
7.1	Data flow	15
8	CLI reference	16
8.1	run — all-in-one synthesis	16
8.2	text-to-wav — audio-only synthesis	16
8.3	tract-to-mp4 — re-render a video	17
8.4	inspect — tract inspector	17
8.5	plot-tract — graphical inspector	17
9	Python API reference	17
9.1	Pipeline — high-level orchestrator	17
9.2	Video re-rendering	18
9.3	Engine and VTL APIs	18

10 Input notation and grapheme-to-phoneme conversion	18
10.1 Engine SAMPA inventory	18
10.2 English text rules	18
11 File formats reference	19
11.1 .tract — tract sequence at 400 Hz	19
11.2 .wav, .mp4, .txt	20
12 Worked example: “this is easy for us”	20
12.1 g2p output	20
12.2 Tract sequence	20
12.3 Audio and video	20
13 Differences with the vtl-pipeline reference repository	21
14 Limitations and known issues	21
15 Developer guide	22
15.1 Package layout	22
15.2 Running the tests	22
15.3 Extending the model	22
16 License, acknowledgments and references	23

1 Introduction

covtl-pipeline is an articulatory text-to-speech and video-synthesis pipeline. Given raw English text, it produces:

- a phonemic transcription in the engine SAMPA notation,
- a tract sequence (**.tract**) giving the 19 VocalTractLab tract parameters and 11 glottis parameters at 400 Hz,
- an audio waveform (**.wav**) at 44 100 Hz, synthesized by VocalTractLab,
- an animated MP4 showing the sagittal cross-section of the vocal tract synchronized with the waveform.

The synthesis engine is the **Berthommier COVTL articulatory model**, a compact mathematical model of speech production in which phonetic knowledge is encoded geometrically:

- every phoneme — vowel or consonant — is a point $z = \rho e^{i\theta}$ in a *polar phonetic plane*; the radius ρ is an aperture/closure magnitude and the angle θ is the direction of the articulatory gesture;
- the plane is projected linearly onto the 15 controlled parameters of a Maeda-type vocal-tract model, so that a single (ρ, θ) target simultaneously coordinates jaw, lips, hyoid, tongue body, tongue tip, tongue base and tongue sides;
- coarticulation is not a post-hoc smoothing: it is the structural *superposition* of a continuous *vocalic branch* and transient *consonantal branch*, each projected onto its own subset of parameters (its *selector*). During a labial stop the tongue keeps following the vowel; during a lingual consonant the jaw and lips keep following it;
- source characteristics (voicing, nasality, laterality, glottal effort, F0) are driven by a tag-based extension engine whose events are *indexed by the articulatory trajectory itself* — the source inherits the coarticulation of the tract.

The model is coupled to **VocalTractLab 2.4** (VTL), the physiological articulatory synthesizer developed by Peter Birkholz and collaborators at TU Dresden. VTL converts each 400 Hz vocal-tract state into cross-sectional areas and runs its aerodynamic-acoustic simulation to produce the 44.1 kHz waveform. The package depends on the **vocaltractlab-cython** wheel (Paul Krug’s Cython binding of VTL), which also provides the SVG export of vocal-tract shapes used by the video pipeline. The repository and packaging layout follow the [vtl-pipeline](#) reference repository; the synthesis path itself is entirely different (see §13).

Key features

- End-to-end pipeline from raw text to MP4 in a single command, or direct SAMPA input with **--phonetic**.
- Polar-coordinate phonetic targets with documented, calibrated coefficient tables (§3).
- Selector-based coarticulation with lightened selectors, contextual consonant targets, cluster scoping and vowel-to-vowel background arcs (§4).
- Tag-based source model: voicing, VOT by place of articulation, nasality with carry-over, glottal effort and F0 declination (§6).

- Syllabic trajectory engine (`syl`) with pause handling and phrase-final holds; a legacy superposition-graph engine remains selectable (§5).
- Inspectors for the tract sequence: `inspect` (tables, ranges) and the graphical `plot-tract`.
- 55 pytest tests: engine non-regression against `regression_baselines.json` plus CLI/g2p unit tests.

2 Quick start

2.1 Prerequisites

- Python 3.9 or newer.
- No system `ffmpeg` is required: the `imageio-ffmpeg` wheel bundles a private `ffmpeg` binary for MP4 encoding.
- About 200 MB of disk for the Python packages.

2.2 Installation

```
git clone https://github.com/fberthommier/covtl-pipeline.git
cd covtl-pipeline
pip install .           # or: pip install -e .   (development)

# Optional: PNG figures of .tract files (plot-tract)
pip install ".[plot]"  # pulls matplotlib

# Verify
vtl-synth --help
```

The install pulls `numpy`, `scipy`, `vocaltractlab-cython`, `Pillow`, `imageio-ffmpeg` and `cmudict` automatically. The console script `vtl-synth` lands in the Python *Scripts* folder; with `pip install --user` on Windows this is `%APPDATA%\Python\Python39\Scripts` — add it to `PATH`, call it by full path, or use the provided `launch.bat`. Under Python 3.9, if `pip` tries to build `vocaltractlab-cython` from source, install `vocaltractlab-cython==0.0.13` instead (full synthesis and video path available; four optional analysis helpers raise a clear error).

2.3 First run

```
vtl-synth run "this is easy for us" -o ./out
```

This produces the following artifacts in `./out/` (real output, v1.0.2 defaults):

<code>this_is_easy_for_us.tract</code>	268,286 bytes	952 frames x (19 tract + 11 glottis)
<code>@ 400 Hz</code>		
<code>this_is_easy_for_us.wav</code>	209,484 bytes	PCM 16-bit, mono, 44100 Hz, 2.37 s
<code>this_is_easy_for_us.mp4</code>	239,001 bytes	H.264 960x840 @ 25 fps + AAC, 2.40 s
<code>this_is_easy_for_us.txt</code>	20 bytes	the SAMPA transcript:
<code>Dis.i.z.i.zi.fOR.@s</code>		

The video shows the sagittal cross-section of the vocal tract; the audio is normalized to 90% of full scale by the pipeline.

2.4 Inspecting the output

```
# Summary + per-parameter min/max/mean of a .tract
vtl-synth inspect ./out/this_is_easy_for_us.tract --states 2 --ranges

# Graphical counterpart: 400 Hz trajectory of each parameter
vtl-synth plot-tract ./out/this_is_easy_for_us.tract --which all -o
./out/tract.png
vtl-synth plot-tract ./out/this_is_easy_for_us.tract --params JX JA TTX TTY

# Re-render a video from an existing tract + wav (no re-synthesis)
vtl-synth tract-to-mp4 ./out/this_is_easy_for_us.tract
./out/this_is_easy_for_us.wav \
-o ./out/this_is_easy_for_us-hd.mp4
```

3 The COVTL model I: phonetic targets in polar coordinates

This section and the next describe the core of the synthesis engine: how phonemes are represented, how the plane is mapped onto vocal-tract parameters, and how coarticulation is obtained. The reference publications are listed in §16; the implementation lives in `vtl_synth/core/polar.py`, `projection.py`, `constants.py` and `syltraj.py`.

The global architecture of the model is a sum of three contributions:

$$\text{VTL}(t) = B(t) + E_{\text{source}}(t) + E_{\text{ortho}}(t), \quad (1)$$

where $B(t)$ is the *Berthommier polar core*: the 15 COVTL tract parameters trajectory built from polar targets; $E_{\text{source}}(t)$ groups the extension parameters produced by the tag/envelope engine (velum VS/VO and the 11 glottis parameters); and $E_{\text{ortho}}(t)$ is the orthogonal branch (TS3 tongue-side shape and the speaker-derived VS), read from the JD3 speaker file (§5.1).

3.1 The (ρ, θ) phonetic plane

Every phoneme is a point

$$z = \rho e^{i\theta}, \quad \rho \geq 0, \theta \in [0, 2\pi), \quad (2)$$

in a complex *phonetic plane*:

- ρ is a dimensionless **aperture/closure magnitude**. $\rho = 0$ is the origin (minimum displacement); $\rho \leq 1$ covers the vowel space; $\rho > 1$ — beyond the unit circle — is the *constriction/occlusion region* used by consonants. The per-parameter amplitude c_0 (Table 3) carries the physical (centimeter) scaling.
- θ is the **direction of the articulatory gesture**. The convention is fixed by the three corner vowels of the Maeda reference vowel triangle: $/a/ = \pi$ (180°, low-open-central), $/i/ = 5\pi/3$ (300°, palatal-front), $/u/ = \pi/3$ (60°, velar-back-round).

Figure 1 shows the plane drawn to scale from the target tables. Two properties are worth noting:

- **Vowels are anchors.** A vowel target becomes a stationary plateau (a *hold*) of the trajectory whenever it is not bracketed by an immediate transition; $\rho \leq 1$ keeps the tract open.
- **Consonants are arc endpoints.** A consonant is never interpolated from the previous consonant — each keeps its own target — and ρ typically ranges from 1.0 to 1.45, i.e. the gesture goes *through* the unit circle to reach a constriction or a complete closure (occlusion areas of 0.000 cm² are reached and verified in the test suite).




Figure 1: The (ρ, θ) phonetic plane of the COVTL model, drawn to scale from the target tables of `constants.py`. Filled dots: the 12 vowel targets (Table 1); crosses: landmark consonant targets (Table 2). The dashed unit circle is the frontier of the vowel space; consonant closures live beyond it ($\rho > 1$). The schwa @ sits at the centroid of the calibrated $a-i-u$ triangle. $/m/$ coincides with $/b/$ — same target and same selector; only the velum extension (VO, outside the COVTL selector) differs.

3.2 Vowel targets

The twelve vowel targets of the engine (VOWEL_TARGETS in `constants.py`) are given in Table 1. The corner vowels $/a/$, $/i/$, $/u/$ are the strict Maeda anchors; the other vowels were fitted by least squares under the constraint $\rho \leq 1$. Two documented compensations: the center coefficient c_1 (TCY) had to be raised by +0.25 to obtain the velar occlusion of $/g/$ (Table 3), and the radii of $/i/$ and $/u/$ were reduced ($0.80 \rightarrow 0.60$ and $0.80 \rightarrow 0.55$) to compensate, preserving F2 to within a few tens of Hz. The schwa @ is placed at the centroid of the calibrated $a-i-u$ triangle, $\bar{\rho} = (1.000 + 0.600 + 0.550)/3 = 0.7167$, $\theta = \pi$; it serves as the mute vowel of isolated consonants and as a waypoint against vowel-to-vowel closure (§4.5).

3.3 Consonant targets and selectors

A consonant target is a triple $(\rho, \theta, \mathcal{S})$: a polar point *plus a selector* $\mathcal{S}$, the list of COVTL parameters that the consonant is allowed to command. Parameters outside $\mathcal{S}$ are *not* touched by the consonant — they keep following the vocalic branch. This is the structural basis of coarticulation (§4).

Two conventions are essential:

- **Voicing is not in the tract.** Only the canonical voiced form of each obstruent is tabulated; the voiceless keys p , t , k , f , s , S , T , tS are normalized to their voiced partners ($p \rightarrow b$,

key	ρ	θ (rad)	θ (deg)	note
a	1.000	π	180.0	strict Maeda anchor
i	0.600	$5\pi/3$	300.0	anchor, ρ compensated
u	0.550	$\pi/3$	60.0	anchor, ρ compensated
e	0.988	4.581	262.5	
E	0.594	3.563	204.1	
o	1.000	1.827	104.7	clamped at $\rho = 1$
0	1.000	2.616	149.9	clamped at $\rho = 1$
6	0.894	3.053	174.9	œ
9	0.722	3.105	177.9	œ̃ (nasalized via VO)
@	0.7167	π	180.0	centroid of the <i>a-i-u</i> triangle
y	0.519	4.548	260.6	
2	0.396	3.946	226.1	∅

Table 1: Vowel targets VOWEL_TARGETS (constants.py).

t→d, k→g, f→v, s→z, S→Z, T→D, tS→dZ) before target lookup, because a voicing pair shares the same place of articulation — hence the same (ρ, θ) and the same selector. The voicing distinction is carried entirely by the glottal source (§6).

- **Glides are consonantal nodes anchored on the reference vowel tree.** j is anchored at the reference coordinates of i (0.80, 300°), w at those of u (0.60, 60°) and ɥ at those of y (0.519, 260.6°); they are treated as C nodes (no plateau), with selectors restricted accordingly. These anchor radii are the reference-model values, deliberately kept independent of the calibrated vowel plateaus of Table 1 (where $\rho_i = 0.60$ and $\rho_u = 0.55$): a glide is a consonantal target, not a vowel plateau.

key	(ρ, θ)	selector	note
b	(1.06, 50°)	JA, LD, LP	labial stop; lightened selector (§4.2)
d	(1.10, 270°)	JA, TCX, TCY, TBX, TBY, TTX, TTY	dental/alveolar stop
g	(1.45, 60°)	JA, TCX, TCY, TBX, TBY	velar stop; conditional → g_vel/g_pal
g_vel	(1.45, 60°)	JA, TCX, TCY, TBX, TBY	velar occlusion 0.000 cm ²
g_pal	(1.30, 300°)	JA, TCY, TBX, TBY	palatal; TCX removed from the selector
v	(1.14, 353°)	JA, JX, LP, LD	labiodental; lightened selector
z	(1.00, 289°)	JA, TCX, TCY, TBX, TBY	alveolar fricative
z_front	(0.90, 289°)	JA, TCX, TCY, TBX, TBY	contextual variant before i, u
Z	(1.02, 351°)	JA, TCX, TBX, TBY, LP	postalveolar; TCY removed (coarticulation)
D	(1.15, 271°)	JA, TCX, TCY, TBX, TBY	dental fricative
R	(1.28, 234°)	JA, TCX, TCY, TBX, TBY, TS2	uvular
dZ	(1.15, 7°)	JA, TCX, TCY, TBX, TBY, LP	affricate (canonical voiced)

key	(ρ, θ)	selector	note
l	(1.15, 7°)	JA, TCX, TCY, TBX, TBV	lateral (laterality via TS3)
L	(1.25, 7°)	JA, TCX, TCY, TBX, TBV	dark lateral
C	(1.15, 353°)	JA, TCX, TCY	palatal fricative (ç)
X	(1.28, 234°)	JA, TCX, TCY, TBX, TBV, TS2	uvular fricative
m	(1.06, 50°)	JA, LD, LP	labial nasal; velum via VO extension
n	(1.10, 7°)	JA, TCX, TCY, TBX, TBV, TTX, TTY	alveolar nasal
N	(1.20, 60°)	(empty)	derived velar place (no closure target)
J	(1.20, 300°)	(empty)	derived dorsal place; the English velar NG maps here (documented compromise, Table 5)
w	(0.60, 60°)	JA, LP, LD, TCX, TCY	glide node on the reference vowel tree
j	(0.80, 300°)	JA, TCX, TCY	glide node on the reference vowel tree
ɥ	(0.519, 260.6°)	JA, TCX, TCY	glide node on the reference vowel tree
h	(0.50, 180°)	(empty)	near-neutral; source-driven

Table 2: Consonant targets `CONSONANT_TARGETS` (`constants.py`), canonical voiced forms. The selector is the set of COVTL parameters commanded by the consonant; all other parameters follow the vocalic branch. Voiceless keys resolve to the voiced partner before lookup.

3.4 Context-conditioned targets

The tabulated target is only the *canonical* form; the engine adjusts it from context, in two ways:

Conditional places. /*g, k*/ choose between `g_vel` and `g_pal` from the vowel context: a following vowel with $\theta > \pi$ (front) triggers the palatal target, otherwise the velar one. Similarly `z_front` replaces `z` before the high vowels `i` and `u` (ρ 1.00 $\rightarrow$ 0.90 to keep the alveolar groove open in the high-tongue context).

Cluster rules. `adjust_consonant_theta` applies the reference place rules inside consonant clusters: `d/t` retracts to the cluster angle $\theta_{CC,d} = -2.5\pi/6$ with $\rho = 1.2$, `g/k` uses $\theta_{CC,g} = -\pi/12$ with $\rho = 1.1$, and the liquids/nasals `l`, `n` are realized stop-like ($\rho_{\text{stop}} = 1.2$, $\theta = 270^\circ$). Outside a cluster, `d/t` use the dental angle $23\pi/16$, and `g/k` back to $\theta_{\text{velar,back}} = 23\pi/12$ after a back vowel.

3.5 Projection onto the VocalTractLab parameters

Level 2 of the model — the *coordination* — maps the phonetic plane onto the 15 COVTL-controlled parameters of a Maeda-type vocal-tract model (the remaining 4 of the 19 VTL tract parameters are extensions: velum VS/VO and the auto-driven tongue root TRX/TRY). For a

param	c_1	c_0	c_2 (rad)	role
HX	0.722233	0.555533	0.000000	hyoid position (x)
HY	-5.155867	1.246761	3.361756	hyoid position (y)
JX	-0.025000	0.050000	2.094395	jaw advance
JA	-3.371067	1.299456	5.354986	jaw aperture
LP	0.382967	0.617041	1.042238	lip protrusion
LD	0.466400	0.583000	4.092457	lip distance
TCX	1.090767	1.528173	5.373564	tongue center (x)
TCY	-0.914133	0.690945	6.031444	tongue center (y)
TTX	4.044400	0.832384	4.466027	tongue tip (x)
TTY	-0.459700	1.084300	0.116283	tongue tip (y)
TBX	2.795833	1.272045	4.907833	tongue base (x)
TBY	0.159033	0.836452	0.076068	tongue base (y)
TS1	0.248667	0.049433	0.545179	tongue shape 1
TS2	0.031067	0.032478	4.484628	tongue shape 2
TS3	0.032933	0.119409	2.897420	tongue shape 3

Table 3: Projection coefficients `CO_VTL` (`constants.py`; tuple order in the code is (c_1, c_0, c_2)). Documented calibration shifts: LD c_1 scaled $\times 0.8$ (avoid negative lip distance at high ρ); TCY $c_1 + 0.25$ (velar occlusion of /g/); TTY $c_1 + 0.15$ (apical occlusion of /d, t/; Maeda origin -0.6097).

static target (ρ, θ) , each parameter p receives

$$P_p(\rho, \theta) = c_{1,p} + \rho c_{0,p} \cos(c_{2,p} - \theta), \quad (3)$$

or, in complex form with $\Psi_p = c_{0,p} e^{ic_{2,p}}$ and $z = \rho e^{i\theta}$,

$$P_p = c_{1,p} + \text{Re}[\Psi_p \bar{z}]. \quad (4)$$

Geometrically: the center c_1 is the parameter value at the origin of the plane; the lever c_0 is the amplitude of the excursion; the phase c_2 is the gesture direction in which parameter p is *most* modified (the maximum of Eq. (3) is reached at $\theta = c_2$). A single θ therefore coordinates all 15 parameters at once — this is how jaw, lips, hyoid and the four tongue degrees of freedom move together from one phonetic direction.

As an example, the jaw aperture during any target (ρ, θ) is $\text{JA} = -3.3711 + 1.2995 \rho \cos(5.3550 - \theta)$: at $\theta = 5.355$ rad ($\approx 307^\circ$) the jaw is maximally open for a given ρ — which is why the palatal-open direction also opens the jaw. The three calibration shifts of Table 3 are the price of coupling the Maeda plane to the JD3 speaker anatomy; they were validated against closure areas (0.000 cm² occlusions) and formant targets.




Figure 2: The two-branch superposition of Eq. (5). Each tract parameter belongs to exactly one branch at each instant: the consonant selector decides which.

4 The COVTL model II: coarticulation

Coarticulation — the fact that successive phonemes overlap and deform each other — is not applied as smoothing on a finished trajectory. It is built into the trajectory generation, at three cooperating levels: *superposition* (who commands which parameter), *arc shape* (how fast the handover between targets happens), and *anchor scoping* (how much of a vowel survives at syllable margins).

4.1 Superposition of the vocalic and consonantal branches

The tract trajectory is the superposition of two complex trajectories evaluated in the phonetic plane:

- the **vocalic branch** $z_v(t)$: a continuous chain of arcs joining vowel anchors (*background*);
- the **consonantal branch** $z_c(t)$: transient excursions from the background towards the consonant targets.

Level 3 of the model composes them parameter by parameter:

$$P_p(t) = S_p^v \text{Re}[\Psi_p \overline{z_v(t)}] + S_p^c \text{Re}[\Psi_p \overline{z_c(t)}], \quad S_p^v + S_p^c = 1, \quad (5)$$

where the selector of the current consonant defines the binary masks: $S_p^c = 1$ iff $p \in \mathcal{S}$, else $S_p^v = 1$. During a consonant, the parameters of its selector follow the consonantal branch while *all the others keep following the vocalic background*. When no consonant is active, Eq. (5) reduces to the pure vocalic projection of the background arc.

4.2 Lightened selectors: the tongue during labials

The selectors were deliberately *shortened* to increase coarticulation, with measured justifications recorded in the source:

- **b (and m)**: TCX/TCY removed — the labial closure only needs the jaw (JA) and the lips (LD, LP). Without tongue parameters in the selector, *the tongue follows the vocalic branch during the labial*; the parasitic velar contact that the raised c_1 (TCY) otherwise produced disappears, with the labial closure itself preserved (0.000 cm², lingual margins 0.126–1.056 cm²).
- **f, v**: same lightening (TCX and TCY removed); the labiodental constriction is commanded by JA, JX, LP, LD only.
- **S, Z**: TCY removed — the postalveolar target projects a tongue center that is too high on the JD3 anatomy; by letting the center follow the vowel, healthy gaps of 0.263–1.230 cm² are obtained over the whole vowel set while the hushing spectrum is preserved.

The partial removal of a single parameter is not sufficient; the lightening was validated parameter pair by parameter.

4.3 The transition arc: an asymmetric S-curve in the plane

Between a departure point $p_d = (\rho_d, \theta_d)$ and an arrival point $p_a = (\rho_a, \theta_a)$, the engine does not interpolate linearly. It sweeps an angular parameter ϑ from $-\pi$ to 0 over D steps and blends the two projections with the weight

$$r(\vartheta) = \cos^{p_{\text{exp}}} \left(\frac{\vartheta}{2} \right), \quad p_{\text{exp}} = 2, \quad (6)$$

which grows from 0 to 1 as an asymmetric S-curve (slow departure, fast arrival). The arrival term carries an extra phase $-(\nu/K)\vartheta$ that curves the path in the phonetic plane:

$$P_p(k) = c_{1,p} + r_k \rho_a c_{0,p} \cos(c_{2,p} - \theta_a - \frac{\nu}{K} \vartheta_k) + (1 - r_k) \rho_d c_{0,p} \cos(c_{2,p} - \theta_d), \quad (7)$$

for $k = 1 \dots D$. This is `arc_B` in `polar.py` (implemented exactly as written; the arrival θ is `np.unwraped` first so that the $\pi \rightarrow -\pi$ wrap-around never produces a spurious 2π jump). The curvature constants distinguish the branches of the reference model: the consonantal arcs use $K_c = 10$, $\nu = 1$, while the vocalic background uses $K_v = 30$ — the vocalic handovers are much straighter. The active `syl` engine runs with $\nu = -1$ and a quasi-straight $K = 1000$ (§5); this is a notational convention only — $\nu = 1$ is recovered when the cosine is written with the opposite sign, as in Eq. (2) of arXiv:2307.02299. The reference values ($K_c = 10$, $K_v = 30$, $\nu = 1$) belong to the deprecated `legacy` engine and the polar-video display.

4.4 Clusters: per-parameter scoping

For a sequence of m consonants between two vowel anchors, the `syl` engine builds a block of $(m + 1)$ sub-arcs of T_{cons} steps each (`build_cluster_pval`). The vocalic background arc between the two anchors is computed first and fills the whole block; then, for each sub-arc k , each parameter's endpoints are set to the consonant target $\mathbf{C}_k$ *iff* the parameter belongs to the selector of $\mathbf{C}_k$ — otherwise the parameter keeps its background value at that frame. The temporal mixing inside each sub-arc uses exactly the reference S-curve of Eq. (6):

$$P[kT + t] = r_t \cdot \text{arr} + (1 - r_t) \cdot \text{dep}, \quad t = 1 \dots T, \quad (8)$$

which guarantees C^0 continuity at every sub-arc boundary (zero jump, verified by test). The active parameter set of the block is the *union* of the selectors of its consonants; a "selector frozen at $\theta = 0$ " never drives a parameter in a palatal context.

4.5 Vowel-to-vowel coarticulation and the schwa waypoint

Anchors are reduced at syllable margins. A vowel never reaches its full target ρ_v when it is on the edge of a syllable: the departure anchor (anticipatory) and completion anchor are scaled ($\delta_o = \delta_e = 0.70$ in the reference; `SYL_COEFCEN` = 0.5 for word-end and syllabic-onset anchors in the `syl` engine), and a word-onset anchor is $[0.5 \rho_v, \theta_v]$. This is why a unstressed vowel between two consonants is *colored* by its neighbors instead of being fully realized.

The background never jumps. Between two anchors the vocalic background is a single continuous arc of $2T_{\text{voy}}$ steps; an earlier conditional jump at vowel onsets was removed precisely because it created discontinuities.

Anti-closure routing through the schwa. A direct polar arc between two distant vowels can close the vocal tract mid-transition (the superposition of the two anchor contributions sweeps through the occlusion region — measured on the `i--u` pair: 0.062 cm^2 at mid-arc over 132 swept vowel pairs). The engine hooks a VTL *closure probe* (minimum lingual tube area at the center of the block, on the 19-parameter resting state); when the probe returns less than the threshold 0.18 cm^2 , the background arc is routed through the schwa @ as two half-arcs with exact continuity at the junctions.

4.6 Coarticulation across word boundaries

The input notation is itself a coarticulation control:

- `.` (and `-`) is the *intra-phrase syllable separator*: the engine coarticulates across it. Within a phrase, orthographic words are joined with `.` after syllabification by onset maximization, so **this is easy** becomes one connected token **Dis.iz.i.zi** — there is *no pause at word boundaries inside a phrase* (the syllabification rule is the model’s answer to “why can *big.bi* be changed to *bi.gbi*”: a single intervocalic consonant attaches to the next onset).
- splitting at `.` obeys the phonotactic rule **C.C** → **split**; **C.V**, **V.C**, **V.V** → **merge** (the dot is ignored between a consonant and a vowel);
- *space* is a short inter-word pause (200 ms default), `|` is the long phrase pause (320 ms default); a comma produces a space and sentence punctuation produces `|`;
- the *schwa principle* caps long runs: a dot is inserted after the first consonant of any 4+-consonant run and after the first vowel of any 3+-vowel run, and a vowel-less word receives a mute `@`.

4.7 Nasal carry-over

Nasality coarticulates *beyond* the nasal segment: velum-opening (VO) events are extended 20 ms past the end of the nasal (**NASAL_CARRYOVER_MS**), and the whole VO trajectory is smoothed with a 20 ms first-order filter — an anti-click margin over VTL’s own 12 ms gesture time constant. Nasalized vowels (9, 6~) receive a partial velum opening of 0.5 (vs 1.0 during nasal consonants).

5 The syllabic trajectory engine (syltraj)

The **syl** engine (**core/syltraj.py**, the default) turns a normalized phoneme string into the 100 Hz polar trajectory $B(t)$. Its contract follows the reference COVTL specification:

1. **Nodes**: every phoneme carries its own polar target; no consonant position is ever interpolated from another consonant; targets are possibly adjusted by the cluster rules (§3.4).
2. **Anchors**: vowels only (held when isolated) plus synthetic *vocalic* anchors: word end ($\text{COEFCEN} \cdot \rho$ of the last vowel), syllabic onset (weighted ρ), word onset ($[0.5, \theta_v]$), and the terminals $[0, 0]$ that bracket the phrase.
3. **Trajectory**: a walk over consecutive anchor pairs (**build_global_pval**), emitting typed blocks (Table 4).

The block schedule is the single time base shared by the tract and the source: the envelope engine consumes **block_info** directly (§6), a deliberate design that guarantees perfect branch synchronization. A phrase-final vowel followed by nothing holds its plateau (*phrase-final hold*); a pause is *not* a transition to a neutral position — the tract holds the decaying anchor.

5.1 The orthogonal branch

Three of the 19 VTL tract parameters are not commanded by the polar model: **TS3** (tongue-side/lateral shape) and the speaker-derived **VS** (velic). They come from the *orthogonal branch* (**core/orthogonal_params.py**), which reads context-conditioned targets from the **JD3.speaker** file: e.g. fricatives push TS3 to +1.0, laterals to −1.0, closures to 0, with per-vowel defaults looked up with the key pattern **<prefix>(<vowel>)** and a (a) fallback. TS3 ramps over 10 ms, VS over 15 ms. The branch is additive on TS3 and priority-based on VS, and can be disabled with **--no-ortho** (the **syl** COVTL core then works alone).

Block kind	Content and duration
plateau	vowel hold, T_{voy} steps
background	vocalic arc between anchors, $2T_{\text{voy}}$ steps
cluster	m consonants, $(m + 1)T_{\text{cons}}$ steps (§4.4)
pause	short (T_{pause}) or long ($T_{\text{pause,long}}$) inter-word pause
decay / attack	fall to / rise from the previous anchor around a pending pause, T_{cons} steps
initial / terminal	phrase bracketing from/to $[0, 0]$

Table 4: Block types emitted by the `syl` engine. Time is counted in gesture steps at the 100 Hz gesture rate (1 step = 10 ms); with the CLI defaults $T_{\text{cons}} = 10$ steps (100 ms), $T_{\text{voy}} = 8$ steps (80 ms, transition arc $2T_{\text{voy}}$), short pause 20 steps (200 ms), long pause 32 steps (320 ms).

5.2 The legacy engine

The `legacy` engine (`--engine legacy`) is the pre-`syltraj` chain: it builds explicit superposition graphs in the complex plane (CV, CVC, CCV graphs per syllable, executed by `execute_superposition_segment`) with the reference constants $K_v = 30$, $K_c = 10$, $\nu = 1$, $p_{\text{exp}} = 2$, scaling the consonantal branch to $n_{\text{arcs}} \cdot T \cdot t_{\text{cons}}$ and the vocalic branch to $2T \cdot t_{\text{voy}}$. It is deprecated (a `DeprecationWarning` is emitted) and kept for non-regression.

6 The source model: envelopes and the glottal source

The source branch produces the velum and glottis extensions $E_{\text{source}}(t)$ of Eq. (1). Its guiding principle, stated in the code: *the source envelope has no timing of its own — it is indexed by the articulatory events already produced by $B(t)$; coarticulation is inherited.*

6.1 The three-level tag system

Phonetic knowledge is described by *tags*, not by raw VTL values, through three levels (`core/tag_system.py`):

Level A — phonetic tags: `voiced`, `nasal`, `continuant`, `strident`, `aspiration`, `place`, `manner`, `rounding`, `height/depth`, `front/back`, `tense`.

Level B — articulatory tags: `glottal_shape` (`modal`, `voiceless`, `breathy`, `whisper`, ...), `velum` (`oral`, `nasal`, `nasalized`), `tongue_mode`.

Level C — control tags: `glottal opening` (0–1), `subglottal pressure` (200–2000 dPa), `F0` (50–500 Hz), `aspiration strength` (0–30 dB), `chink area` (0–50).

A tag like `<glottal = voiceless>` is preferred to a hard-coded value like `V0 = 0.42`: phonetic tags are compiled to speaker parameters at run time, which is what makes speaker recalibration (JD2/JD3) possible.

6.2 The amplitude envelope $E(t)$ and the source timetable

An *effort* (amplitude) envelope $E(t)$ is built from a token reduction of the phrase (vowels V, plosives C/c, fricatives R/r, nasals N, laterals L, silences O, final F, glide onsets y) with per-token minimum amplitudes (V 1.0, nasals and liquids 0.3, glides 0.5). The transitions between tokens use raised-cosine/Hanning primitives — e.g. a vowel-to-vowel junction stays at 1, a vowel-to-plosive falls as `hann`³, a plosive-to-vowel rises as `hann`. Mapped onto the `syl` block schedule, plateaus and backgrounds get $E = 1$, pauses and terminals $E = 0$, clusters a reduced consonantal

envelope, and decay/attack blocks the fading edges. The envelope is computed at 10 kHz and block-averaged down to 100 Hz.

On top of $E(t)$, a **TimedEnvelope** schedules *events* anchored to the articulatory labels (the vowel triplets (V_o, V, V_e) , the consonant labels $\mathbf{C_p}$, ...) with raised-cosine ramps. The ramp constants (in ms): attack/release 30/30, voicing 40, velum 15, aspiration 8, F0 15, effort 20.

6.3 Voicing, VOT and nasality

Voicing $V(t)$ is built per syllable from the (V_o, V, V_e) triplets: voicing starts at the first voiced consonant of the syllable (the synthetic onset anchor V_o is purely articulatory — silent), stops 10 ms before a stop coda, ramps 30 ms in / 40 ms out; voiceless events then override inside obstruents. During cluster blocks, voiced consonants (stops, nasals, liquids, voiced fricatives) carry $V = 1$ — the pre-voicing bar of /b, d, g/.

Voice-onset time is set by place of articulation (Keating 1984; Lisker & Abramson 1964), in ms:

place	labio-dental	labial	dental	alveolar	post-alveolar	palatal	uvular	velar	glottal
VOT (ms)	25	20	30	40	40	35	45	50	60

Nasality is the velum opening VO of §4.7: 1.0 during nasal consonants, 0.5 on nasalized vowels, baseline -0.1 , with the 20 ms carry-over and smoothing.

6.4 Glottal source, pressure and effort

The 11 glottis parameters interpolate between **presets extracted from the JD3.speaker file** (**modal** — the default, f0 102.216 Hz, pressure 7850 dPa; **stop**; **voiced-fricative**; **voiceless-fricative**; **whisper**; ...), blended by $V(t)$ and smoothed with a second-order critically damped Butterworth filter (time constant 12 ms). The subglottal pressure follows the effort:

$$P_{\text{sub}} = P_{\text{min}} + \min(E, 1.5)^{1.2} (P_{\text{max}} - P_{\text{min}}), \quad P_{\text{max}} = 8000 \text{ dPa}, \quad (9)$$

with a phonation floor of 3000 dPa while voicing is on and a pre-voicing regime during voiced closures (pressure 1800 dPa, relative amplitude clipped to 0.10–0.18). The pulmonary-effort variant zeroes the effort on plosive occlusions (5 ms smoothing) and takes the maximum of the effort, 60 % of the pre-voicing level and 50 % of nasality before smoothing (10 ms).

A per-vowel effort gain compensates the intrinsic intensity deficit of high vowels: **i**, **u**, **e** $\times 1.5$, **y** $\times 1.3$, **2** $\times 1.2$ on their plateaus ($\approx +3$ dB).

6.5 F0: declination and micro-prosody

The default F0 is the JD3 speaker default 102.216 Hz. Each phrase receives a French-style **declination** (Jun & Fougeron 2002): $+3\%$ at onset, -15% at the end, linear in between; phrases are detected as voiced spans separated by more than 200 ms of silence. Superimposed on it, a **micro-prosody** couples F0 to the effort envelope:

$$F0(t) = F0_{\text{base}} + k_E (E(t) - \overline{E}), \quad k_E = 10 \text{ Hz}, \quad (10)$$

clamped to $[40, 600]$ Hz, with the last voiced value held through unvoiced spans.

7 Architecture and data flow

The package is organized into five subpackages:

vtl_synth.core the synthesis engine: notation/parsing, phoneme tables, the polar model, the trajectory engines, the envelope/source modules, and the Pipeline orchestrator (26 engine modules).


vtl_synth.cli the `vtl-synth` entry point with the subcommands `run`, `text-to-wav`, `tract-to-mp4`, `inspect`, `plot-tract`.

vtl_synth.vtl the VocalTractLab API facade (`api.py`) over the `vocaltractlab-cython` wheel.

vtl_synth.video the sagittal renderer (`renderer.py`: SVG frames $\rightarrow$ PNG), the MP4 encoder (`encoder.py`: ffmpeg H.264/AAC), and the parameter figure tool (`tract_figure.py`).

vtl_synth.utils the grapheme-to-phoneme module (`g2p.py`, CMUdict).

7.1 Data flow



Each stage of the chain:

1. **g2p**: English text $\rightarrow$ connected engine SAMPA tokens (§10); or direct SAMPA with `--phonetic`.
2. **Parse & normalize**: Unicode/IPA normalization, longest-match tokenization, phrase splitting on |, syllable structure.
3. **B(t)**: the `syl` engine builds the 100 Hz polar trajectory block by block (§5), projected through Eq. (3) into the 15 COVTL parameters.
4. **Source**: the envelope $E(t)$, voicing $V(t)$, velum VO/VS, glottis trajectory and F0 (§6), all indexed by the block schedule; the orthogonal branch adds TS3/VS (§5.1).
5. **Assemble**: the 15 COVTL parameters are scattered into the 19-parameter VTL vector (VS at index 6, VO at 7; TS3 additive; TRX/TRY left at 0 — VTL drives them), the glottis block is appended, and everything is linearly up-sampled 100 $\rightarrow$ 400 Hz into the `.tract` file.
6. **VTL synthesis**: each tract state $\rightarrow$ area function (40 tube sections, JD3 anatomy) $\rightarrow$ aerodynamic-acoustic simulation $\rightarrow$ 110 audio samples per state at 44.1 kHz.
7. **Video**: every video frame is exported as an SVG of the sagittal cross-section by VTL itself, rasterized to PNG, and encoded with the bundled `ffmpeg`.

Three sampling rates

- **Model rate:** 100 Hz (gesture step 10 ms) — the polar trajectory and the envelopes.
- **Tract rate:** 400 Hz (`.tract`) — 19 tract + 11 glottis parameters per frame, 2.5 ms hop.
- **Audio rate:** 44 100 Hz — 110 audio samples per tract state (`.wav`, PCM 16-bit mono).

8 CLI reference

The `vtl-synth` command exposes five subcommands. The engine options (COVTL model parameters) are shared by `run` and `text-to-wav`:

Engine flag	Description
<code>--phonetic</code>	input is already engine SAMPA (skip g2p)
<code>--tcons MS</code>	consonant (sub-arc) duration, default 100 ms
<code>--tvoy MS</code>	vowel gesture duration, default 80 ms
<code>--tpause MS</code>	short inter-word pause (space), default 200 ms
<code>--tpause-long MS</code>	long phrase pause (l), default 320 ms
<code>-f0, --f0 HZ</code>	fundamental frequency, default 102.216 Hz (JD3)
<code>-s, --speaker</code>	speaker name, default JD3
<code>--no-ortho</code>	disable the orthogonal branch (TS3/VS)
<code>--engine {syl,legacy}</code>	trajectory engine, default <code>syl</code>

8.1 run — all-in-one synthesis

```
vtl-synth run "this is easy for us" [options]
```

Runs the full pipeline: text → `.tract` → `.wav` → `.mp4` (+ the SAMPA transcript `.txt`).

Flag	Description
<code>-o, --output-dir</code>	output directory (default: <code>./out</code>)
<code>--label</code>	filename stem (default: derived from the text)
<code>--no-video</code>	skip MP4 rendering
<code>--fps</code>	video frame rate (default: 25)
<code>--scale</code>	pixels per SVG unit (default: 2 → 960×840)

```
# Direct phonetic input (no g2p): syllables '.', short pause ' ', long '/'
vtl-synth run --phonetic "ba da ga | iowa a g" -o ./out

# Slower, more coarticulated speech
vtl-synth run "this is easy for us" -o ./out --tcons 140 --tvoy 110

# Audio only
vtl-synth run "hello world" -o ./out --no-video
```

8.2 text-to-wav — audio-only synthesis

```
vtl-synth text-to-wav "hello world" -o ./out
```

Same engine path as `run`, stops after the `.wav`.

8.3 tract-to-mp4 — re-render a video

```
vtl-synth tract-to-mp4 <tract> [<wav>] [-o output.mp4] [--fps 25] [--scale 2]
```

Re-renders the sagittal video from an existing `.tract`; when the `.wav` is omitted it is re-synthesized from the tract. Useful to change resolution or frame rate without re-running the model.

8.4 inspect — tract inspector

```
vtl-synth inspect ./out/this_is_easy_for_us.tract --states 2 --ranges
```

Prints the shape, rate and duration of the tract sequence, the first/last N states as per-parameter tables, and with `--ranges` the min/max/mean of each of the 19 tract parameters.

8.5 plot-tract — graphical inspector

```
vtl-synth plot-tract ./out/this_is_easy_for_us.tract --which all -o
./out/tract.png
vtl-synth plot-tract ./out/this_is_easy_for_us.tract --params JX JA TTX TTY
--ranges
```

Draws the 400 Hz trajectory of every tract parameter into a PNG figure (shared time axis, one subplot per parameter); `--which all` adds the glottis `f0` and pressure, `--params` restricts the tract block, `--ranges` prints the inspect table. Requires the optional plot extra (`pip install "[plot]"`).

9 Python API reference

9.1 Pipeline — high-level orchestrator

```
from vtl_synth import Pipeline

pipe = Pipeline(
    f0_hz=102.216,          # JD3 default; per-phrase declination applied
    speaker="JD3",
    engine="syl",           # 'syl' (recommended) or 'legacy'
    use_orthogonal=True,    # TS3/VS orthogonal branch
    t_cons_ms=100.0,        # consonant sub-arc duration
    t_voy_ms=80.0,          # vowel gesture duration
    pause_short_ms=200.0,   # inter-word pause
    pause_long_ms=320.0,    # end-of-sentence pause
)

result = pipe.run("this is easy for us", output_dir="./out")
print(result.sampa)        # 'Dis.iz.i.zi.fOR.@s'
print(result.tract_path, result.n_frames, result.duration_s)

# phonetic input, audio only
result = pipe.text_to_wav("ba da ga | iowa a g",
                          output_dir="./out/phon", use_g2p=False)
```

The `PipelineResult` dataclass exposes `sampa`, `phrases`, the artifact paths (`tract_path`, `wav_path`, `mp4_path`), `n_frames` (at 400 Hz), `duration_s` and `warnings`.

9.2 Video re-rendering

```
info = Pipeline.tract_to_mp4(
    tract_path="./out/this_is_easy_for_us.tract",
    wav_path="./out/this_is_easy_for_us.wav",
    out_path="./out/regenerated.mp4",
    fps=25, scale=2,
)
```

This does *not* re-run the model: it parses the tract sequence, renders one SVG frame per video frame and encodes with ffmpeg.

9.3 Engine and VTL APIs

The full engine API (63 public symbols) is available from `vtl_synth.core`: the orchestrators (`Pipeline`, `build_phrase_tract`, `text_to_tract`), the polar model (`compute_P`, `arc_B`, `polar_arc`, `project_covtl`, `superimpose`), the trajectory engines (`build_phrase_pval`, `build_global_pval`), the envelope/source modules (`TimedEnvelope`, `FeatureTimer`, `build_glottis_trajectory`), the constants (`VOWEL_TARGETS`, `CONSONANT_TARGETS`, `CO_VTL`, `CO`) and the assembly layer (`AssembleTract`, `write_tract_file`).

```
from vtl_synth.vtl.api import (synthesize_to_wav, render_tract_svg,
                               get_vtl_info, get_vowel_shape,
                               get_transfer_function)

info = get_vtl_info()
# {'version': ..., 'speaker': 'JD3', 'sr_audio': 44100,
#  'sr_internal': 400.0, 'n_tract_params': 19,
#  'n_glottis_params': 11, 'n_samples_per_state': 110,
#  'n_tube_sections': 40, ...}

synthesize_to_wav("./out/this_is_easy_for_us.tract",
                  "./out/resynthesized.wav")
```

10 Input notation and grapheme-to-phoneme conversion

10.1 Engine SAMPA inventory

- **Vowels:** a e i o u y E ʊ @ 2 6 9 (12 targets, Table 1);
- **Consonants:** p b t d k g f v s z S Z T D tS dZ m n J l L R j w h ʉ C X (canonical voiced forms of Table 2 plus their voiceless partners; R is the only rhotic — a spelled r in the input is normalized to R by `notation.py`);
- **Separators:** . syllable break (coarticulated), space short pause, | long pause; ~ marks nasality.

10.2 English text rules

With English text input (the default), the bundled g2p applies:

1. CMUdict lookup of each word (first pronunciation), stress digits stripped;
2. ARPAbet → engine SAMPA mapping (Table 5);
3. syllabification by onset maximization: the vowel closes the nucleus, a single intervocalic consonant attaches to the next onset, longer clusters split before the last consonant, word-final consonants form the coda;

ARPAbet	SAMPA	Note
P B T D K G	p b t d k g	direct (voiceless resolves to voiced target)
CH, JH	tS, dZ	affricates
F V S Z SH ZH	f v s z S Z	direct
TH, DH	T, D	dental fricatives — engine targets exist
HH	h	direct
M N NG	m n J	nasals; NG→J is a documented compromise: the English velar nasal [ŋ] is mapped to the derived dorsal node (neither N nor J carries a closure target, Table 2)
L R W Y	l R w j	R is the uvular of the engine inventory
AA AE	a	AE merges to a
AH	@	schwa
AO	O	
EH EY	E, e	
IH IY	i, i	quantity not distinguished
OW	o	
UH UW	u, u	quantity not distinguished
AW AY OY ER	a w, a j, O j, @ R	diphthongs/rhotic → key sequences

Table 5: ARPAbet → engine SAMPA mapping (ARPA_TO_SAMPA, core/phonemes.py).

- the words of a phrase are joined with . into connected tokens — the engine coarticulates across word boundaries (§4.6).

Punctuation: sentence-ending punctuation (.!?:;) becomes the long phrase pause |; a comma or dash becomes the short inter-word pause. Unknown words that are already valid engine SAMPA (`is_sampa_token`) are passed through verbatim; other unknown words are kept with a warning.

11 File formats reference

11.1 .tract — tract sequence at 400 Hz

Plain-text file: comment lines starting with #, a fixed header (sample rate, parameter counts and names, frame count), then one line per frame with the 19 tract + 11 glottis values (%.6f):

```
# VTL tract sequence file
# version: 1
# sample_rate: 400
# n_tract_params: 19
# n_glottis_params: 11
# n_frames: 952
# tract_param_names: HX HY JX JA LP LD VS VO TCX TCY TTX TTY TBX TBY
#                    TRX TRY TS1 TS2 TS3
# glottis_param_names: f0 pressure x_bottom x_top chink_area lag rel_amp
#                    double_pulsing pulse_skewness f0_flutter
#                    aspiration_strength
# ---
<frame 0: 19 tract values, then 11 glottis values>
...
```

The 19 tract parameters: hyoid HX HY, jaw JX JA, lips LP LD, velopharyngeal port VS VO (extensions), tongue center TCX TCY, tongue tip TTX TTY, tongue base TBX TBY, tongue root TRX TRY (auto-driven, kept at 0), tongue shape TS1 TS2 TS3. The 11 glottis parameters: f0 (Hz),

pressure (dPa), x_bottom, x_top, chink_area (posterior opening/breathiness), lag, rel_amp, double_pulsing, pulse_skewness, f0_flutter (%), aspiration_strength (dB).

The reader is `read_tract` (`vtl_synth.video.renderer`); it parses `# sample_rate:` and expects at least 19 columns.

11.2 .wav, .mp4, .txt

- `.wav`: PCM 16-bit mono at 44 100 Hz, peak-normalized to 0.9 full scale by the pipeline.
- `.mp4`: H.264 (yuv420p) + AAC; default 960×840 at 25 fps (`--scale 2`). The upper panel is the sagittal cross-section exported by VTL as SVG; the audio track is muxed by ffmpeg (`-shortest`).
- `.txt`: the engine SAMPA transcript actually synthesized (e.g. `Dis.iz.i.zi.fOR.@s`) — the audit trail of the g2p.

12 Worked example: “this is easy for us”

This section walks through the real v1.0.2 output of the demo phrase (all numbers measured on the artifacts of a default run).

12.1 g2p output

word	ARPAbet (CMUdict)	engine token
this	DH IH1 S	Dis
is	IH1 Z	iz
easy	IY1 Z IY0	i.zi
for	F AO1 R	fOR
us	AH1 S	@s

The words of the phrase are connected with '.', giving the single engine token `Dis.iz.i.zi.fOR.@s`. Note the resyllabifications (*is easy* → *iz.i.*, *easy* split *i.zi*): one coarticulated phrase, no pause at the orthographic word boundaries. The full SAMPA transcript written to `.txt`: `Dis.iz.i.zi.fOR.@s`.

12.2 Tract sequence

```
shape      : (952, 30)  (11 glottis cols)
rate       : 400 Hz
duration   : 2.380 s
```

952 frames at 2.5 ms hop. The first and last frames are the resting state — the projection of the phrase-terminal anchors [0, 0] is the center line of Table 3, e.g. `HX = 0.722`, `JA = -3.371`, `TCX = 1.091`.

12.3 Audio and video

```
wav : PCM 16-bit, 1 channel, 44100 Hz, 104720 frames, 2.3746 s
      peak 90.00 % FS (pipeline normalization), RMS 7.6 % FS
mp4  : h264 (High), yuv420p, 960x840, 25 fps, 2.40 s
      + aac 44100 Hz mono; 239,001 bytes
```

The video and audio durations agree to within one video frame (2.40 vs 2.37 s), confirming the sync (each frame decimates the 400 Hz tract sequence by $400/25 = 16$ frames).

13 Differences with the vtl-pipeline reference repository

The [vtl-pipeline](#) reference repository chains VTL-native stages (`text` $\rightarrow$ `.seg` $\rightarrow$ `.ges` $\rightarrow$ `tract` $\rightarrow$ `wav/mp4`) through a ctypes binding to `libVocalTractLabApi`. `covtl-pipeline` shares its packaging and CLI layout but synthesizes the `tract` sequence **directly** with the Berthommier COVTL model:

	vtl-pipeline (reference)	covtl-pipeline
Tract generation	VTL segment $\rightarrow$ gestural score	COVTL polar model, no <code>.seg/.ges</code>
Coarticulation	VTL gesture blending	selector superposition (§4)
Targets	VTL shape library	(ρ, θ) plane (§3)
VTL backend	ctypes <code>.so</code> + Krug wheel	<code>vocaltractlab-cython</code> wheel
Input	English text	English text or SAMPA (<code>--phonetic</code>)

Consequences: no `.seg/.ges` files exist in the output (`inspect` rejects them with an explanatory message), no manual DLL/library management is needed, and the articulatory trajectory is fully transparent — every 100 Hz model sample derives from the documented equations of §3–§4.

14 Limitations and known issues

- **English phoneme approximations.** The engine works on a French-style SAMPA inventory: R is uvular (English rhotics are approximated), **AE/AH** merge to **a/ə**, quantity contrasts (**IH/IY**, **UH/UW**) are not distinguished, and **ER** becomes the sequence **ə R**. Dental fricatives **TH/DH** do have engine targets (T/D at 271°).
- **Fixed gesture durations.** One `--tcons/--tvoy` pair applies to the whole utterance; there is no per-phoneme duration model, no stress-weighted durations and no word-final lengthening beyond the phrase-final hold.
- **Derived places.** N and J have empty selectors (derived velar/dorsal places); **h** is a near-neutral target driven by the source.
- **Legacy engine deprecated.** `--engine legacy` emits a `DeprecationWarning` and is kept only for non-regression.
- **Video renderer.** The `SVG` $\rightarrow$ `PNG` rasterizer is a lightweight polyline renderer built on Pillow (no cairo dependency); complex SVG features beyond the VTL sagittal output are not supported.
- **Python 3.9 compatibility mode.** With `vocaltractlab-cython==0.0.13`, the four analysis helpers added in 0.0.16+ (`get_cross_sections`, `get_centerline`, `get_outlines`, `get_active_speaker`) raise a clear error; synthesis and video are unaffected.

15 Developer guide

15.1 Package layout

```

covtl-pipeline/
|-- pyproject.toml          PEP 621 metadata (Python >= 3.9), GPL-3.0-or-later
|-- LICENSE                GPL-3.0 full text
|-- THIRD_PARTY_NOTICES.md upstream author credits
|-- README.md, INSTALL.md, CHANGELOG.md
|-- launch.bat             Windows quick-launch script
|-- regression_baselines.json E2E non-regression baselines
|-- examples/demo.py       minimal usage example
|-- tests/                 55 pytest tests (9 modules)
`-- vtl_synth/
    |-- __init__.py         public API exports (Pipeline, ...)
    |-- core/              26 engine modules (model, trajectories, source)
    |-- cli/main.py         vtl-synth entry point
    |-- video/             renderer, encoder, tract_figure
    |-- vtl/api.py         VocalTractLab API facade
    |-- utils/g2p.py        CMUdict g2p
    `-- data/vtl_binaries/JD3.speaker

```

15.2 Running the tests

```

pip install -e ".[dev]"
pytest -v                                # all 55 tests
pytest tests/test_polar_model.py
pytest -k "syltraj"

```

The suite covers the polar model (projections, arcs), the `legacy/syl` equivalence and structure tests, glottal classification, notation/parsing, CLI and g2p, the tract figure, and end-to-end non-regression against `regression_baselines.json`.

15.3 Extending the model

- **New phoneme target:** add the (ρ, θ) entry (and selector) to `CONSONANT_TARGETS` or to `VOWEL_TARGETS` in `core/constants.py` — the tokenizers pick it up from `ALL_PHONEME_KEYS` automatically.
- **New pause/gesture timing:** the CLI flags map to Pipeline arguments; the block schedule constants of the `syl` engine are documented in `syltraj.py`.
- **Coding conventions:** see `docs/CONVENTIONS.md` (header format, module docstrings, validation style).
- **SPDX headers** (`SPDX-FileCopyrightText`, `SPDX-License-Identifier`) are present in every `.py` file (GPL-3.0-or-later) and must be preserved in contributions; `THIRD_PARTY_NOTICES.md` must be updated whenever a new upstream dependency is added.

16 License, acknowledgments and references

covtl-pipeline is distributed under the [GNU General Public License v3 or later](#) (GPL-3.0-or-later), Copyright (C) 2026 Frédéric Berthommier. The GPL is required because the package depends on GPL-3.0 upstream components:

- **VocalTractLab** — Peter Birkholz, TU Dresden. <https://www.vocaltractlab.de>
- **vocaltractlab-cython** — Paul Krug. <https://github.com/paul-krug/VocalTractLab-Python>
- concepts from **create_vtl_corpus** — Konstantin Sering, Universität Tübingen. https://github.com/quantling/create_vtl_corpus

The JD3.speaker file is redistributed as package data (`vtl_synth/data/vtl_binaries/`); cmudict is BSD-2. When redistributing, keep the full LICENSE and THIRD_PARTY_NOTICES.md intact and publish modifications under GPL-3.0-or-later.

References

If you use covtl-pipeline in academic work, please cite the publications of the Berthommier COVTL model:

1. F. Berthommier, “A mathematical model of the vowel space”, arXiv:2111.00868 [eess.AS], 2021. DOI: 10.48550/arXiv.2111.00868.
2. F. Berthommier, “Why can big.bi be changed to bi.gbi? A mathematical model of syllabification and articulatory synthesis”, arXiv:2307.02299 [eess.AS], 2023. DOI: 10.48550/arXiv.2307.02299.
3. F. Berthommier, “Synthèse de syllabes avec un modèle de Maeda piloté par une représentation complexe”, in *Actes des 35èmes Journées d’Études sur la Parole (JEP)*, pp. 541–550, ATALA and AFPC, 2024. <https://aclanthology.org/2024.jeptalnrecital-jep.55/>
4. F. Berthommier, “Articulatory modeling of the S-shaped F2 trajectories observed in Öhman’s spectrographic analysis of VCV syllables”, *Interspeech 2025*, Rotterdam, Netherlands. HAL: hal-05233039. <https://hal.science/hal-05233039v1>

Please also cite VocalTractLab itself (see THIRD_PARTY_NOTICES.md for the up-to-date list).